Danish Siddique

Senior Product Engineer ▪ AI Products

idanishsiddique@gmail.com · +92 311 5238087
linkedin.com/in/danishs360 · github.com/danishs360
danishsiddique.com · Islamabad, Pakistan · Remote

Senior product engineer with 6 years of turning real user problems into AI products that run in production. Comfortable across the whole lifecycle: shaping the problem with product and design, then building the agents, backend, apps and infrastructure, and owning reliability and cost after launch. At Kindred AI, works across products used by **1,000+ therapists**, **8,300+ people at 21 companies** and **25,000+ Stellantis employees**, with a focus on the gap between an AI demo and a product people trust: evals, memory, guardrails, latency and cost.

EXPERIENCE

Kindred AI / Senior Product Engineer

Sep 2023 – Present

Remote · AI products for mental-health care and workplace wellbeing, built on a shared AI platform · kindredeq.com

Kindred Workplace · Emotional-intelligence platform with Paige, an AI coach, used by **8,300+ people across 21 companies** · kindredworkplace.com

- Built the **multi-agent backend** behind Paige, the AI coach: a supervisor that routes each message to specialist agents, streamed replies, and an eval suite of golden conversations that alerts the team when quality drops.
- Shipped **Analyze**, which turns recorded conversations of up to 90 minutes into coaching feedback, and **Voice ID**, which recognises the member's own voice, with per-user usage limits to keep AI cost predictable and consent and GDPR deletion built in.
- Took Paige beyond the web app into **Slack and Microsoft Teams** with scheduled check-ins, and connected the ElevenLabs voice agent to live product data securely.

Kindred Practice · HIPAA practice management and telehealth · kindredpractice.com

- Product and engineering owner, working with product, design and clinical stakeholders as the platform grew to **1,000+ therapists**: scheduling, billing, clinical notes, live video sessions (LiveKit, AWS Chime, Deepgram) and AI features.
- Took RAG from prototype to production with eval-driven iteration, cutting **p95 latency ~40%** and **LLM spend ~30%** through prompt caching and model routing.
- Cut the monthly AWS production bill by **61%** by right-sizing instances and removing idle, duplicated and over-provisioned resources.

Agent Ops · Internal platform where AI agents take on engineering and support work

- Built coding agents that turn Sentry errors and tickets into ready-to-review pull requests, **30+ a month**, with a human approving each plan and a review agent checking every PR.
- Built a support agent that drafts customer replies and hands off to the team when unsure, a pipeline that keeps help articles up to date, and a **scanner that blocks patient data and secrets** from agent changes.
- Wrote the shared Claude Code skills and Cursor rules the engineering team builds with, cutting new-hire ramp from weeks to days.

Kindred Core · Multi-tenant AI platform shared by every Kindred product and partner

- Built the **long-term memory** that lets AI agents remember what users shared and when it was true, with full erasure on request. **Roughly doubled accuracy** on time-sensitive questions.
- Tuned speaker recognition on real production audio so it identifies the right person across different microphones, reaching **24 of 24 correct decisions** in testing.
- Built usage metering, billing credits and an audit log of every AI decision for partner companies, with per-partner data rules and complete deletion of a person's records on request.
- Built the **evaluation system** that tests AI behaviour before every release: automated checks that can block a bad release, plus model-graded quality reviews across 24 simulated user personas.

Client platforms

- Medimap (4,200+ walk-in clinics across six provinces): rebuilt the platform after acquisition on Next.js and Payload CMS and launched online booking. **Revenue and users each grew 15%** in two months.
- Stellantis Skills Platform (25,000+ employees, 14 languages): delivered multi-rater skill assessments, SAML single sign-on and manager-based permissions, and set up AWS production with idle servers sleeping to cut cost.

Extendly / Full-Stack Developer

Apr 2023 – May 2024

Remote · Services and integrations for agencies on GoHighLevel · extendly.com

- Built a retrieval-backed support assistant on OpenAI, FastAPI and PostgreSQL that answered **500–600 tickets a day** from the product knowledge base, freeing the support team for cases that needed judgment.

- Built the customer-facing React apps and a serverless AWS Lambda layer that syncs GoHighLevel with third-party tools, so client workflows ran without anyone copying records by hand.

Zaytrics / Full-Stack Developer

Sep 2020 – Apr 2023

Islamabad · Software and AI services · zaytrics.com

- Built a computer-vision system that detects trademark infringement in logos at **85%+ accuracy**, fed by a Selenium data pipeline, removing roughly **70% of the legal team's manual review**.
- Redesigned PostgreSQL and MongoDB schemas that were slowing down as data grew, restoring query performance and leaving room to scale.

INDEPENDENT PRODUCTS

Built and run outside full-time work

TallyReel · AI video creation and social scheduling · tallyreel.com

- Turns one line into a finished short video (stories, explainers, animated and 3D styles, AI presenters, quizzes, ads) and schedules it to **15 social networks**, across web (Next.js), iOS and Android (Expo) and a Mac and Windows desktop app.
- Built a **remote MCP server with 52 tools** and OAuth so Claude and ChatGPT users can create, schedule and review videos, with the same permissions and plan limits as the app.
- Built an in-app AI agent that plans and schedules content, with three autonomy levels so users choose how much it does alone, and per-video GPU cost tracking on Modal with daily model budgets.

Round Zero · AI interview practice · joinroundzero.com

- Live AI voice interviewer on LiveKit that adapts follow-up questions in real time, scores answers against role-specific rubrics calibrated by seniority, and runs coding rounds.

Smaller tools

- **Blurt**: hold a key and talk, and clean text appears wherever you type, with on-device speech recognition in Rust for Mac and Windows.
- **Maina**: voice cloning from 30 seconds of speech in English, Hindi and Urdu, with a built-in AI watermark and evals for word error rate and speaker similarity.

SKILLS

Languages	TypeScript, JavaScript, Python, Rust, SQL
AI and LLMs	Claude, OpenAI, Gemini, DeepSeek, Claude Agent SDK, LangChain, MCP, RAG, graph RAG, multi-agent systems, tool calling, agent memory, LLM evals, Langfuse
Voice and ML	LiveKit, ElevenLabs, Deepgram, Whisper, speaker recognition, computer vision, MLX, llama.cpp, ONNX
Frontend and mobile	React, Next.js, Expo, React Native, Tailwind CSS, TanStack Query
Backend and data	Node.js, NestJS, Fastify, FastAPI, Django, PostgreSQL, pgvector, MongoDB, DynamoDB, Redis, Elasticsearch, Inngest
Cloud and DevOps	AWS, Azure, GCP, Cloudflare, Modal, Docker, Pulumi, GitHub Actions, CI/CD, Sentry, Jest
Auth and billing	SAML SSO, WorkOS, Clerk, RBAC, multi-tenancy, Stripe, RevenueCat, HIPAA, GDPR
AI dev tools	Claude Code, Cursor, Codex, GitHub Copilot

EDUCATION

B.S. Computer Science / Air University, Islamabad

2017 – 2021

- **Final-year research:** deep-learning model that segments the liver and liver tumors in abdominal CT scans, an early start in applied machine learning on medical data.
- **Core coursework:** Data Structures and Algorithms, Database Systems, Operating Systems, Artificial Intelligence, Software Engineering.